



Advanced Database Project

Final Phase

Team 16

Esraa Mamdouh

Khaled Sabry

Mohamed Emad

Roba Gamal

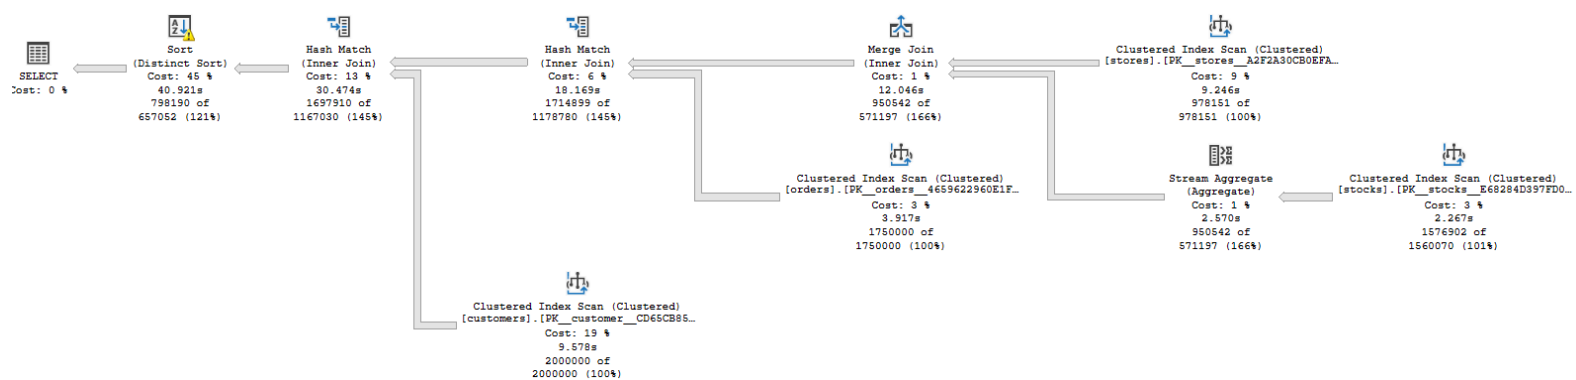
Submitted to:

E. Mostafa Mahmoud

SQL Query #1

```
SELECT DISTINCT
    c.first_name,
    c.last_name
FROM
    sales.orders AS o,
    sales.customers AS c
WHERE
    o.customer_id = c.customer_id
    AND o.store_id IN (
        SELECT
            sto.store_id
        FROM
            sales.stores AS sto
    WHERE
        sto.store_id IN ( SELECT stc.store_id FROM production.stocks AS stc
                        WHERE stc.quantity > 1000 ))
```

Execution Plan Diagram



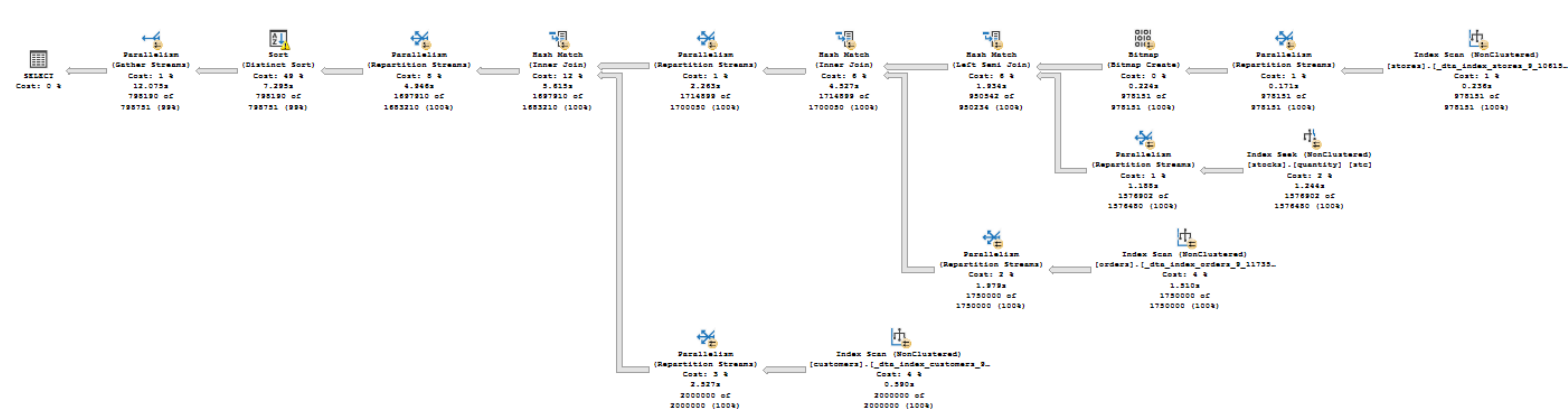
Some Analysis

Rows	StmtText	NodeId	Parent	PhysicalOp	LogicalOp	Argument	TotalSubtreeCost	Parallel
798190	SELECT DISTINCT c.first_name,c.last_nameFROMSales.orders AS o ,sales.customers AS c	1	0				238.8167	False
798190	--Sort(DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC))	2	1	Sort	Distinct Sort	DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC)	238.8167	False
1697910	--Hash Match(Inner Join, HASH:([o].[customer_id])=[c].[customer_id]), RESIDUAL([4	2	Hash Match	Inner Join	HASH:([o].[customer_id])=[c].[customer_id]), RESIDUAL([stor	130.845	False
1714899	--Hash Match(Inner Join, HASH:([stc].[store_id])=[o].[store_id])	6	4	Hash Match	Inner Join	HASH:([stc].[store_id])=[o].[store_id])	55.42873	False
950542	--Merge Join(Inner Join, MERGE:([sto].[store_id])=[(stc].[store_id]), RESIDUAL	7	6	Merge Join	Inner Join	MERGE:([sto].[store_id])=[(stc].[store_id]), RESIDUAL:([store],[33.39322	False
978151	--Clustered Index Scan(OBJECT:([store].[sales].[stores].[PK_stores_A2F2	8	7	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[stores].[PK_stores_A2F2A30C0BEFA	20.97406	False
950542	--Stream Aggregate(GROUP BY:([stc].[store_id]))	9	7	Stream Aggregate	Aggregate	GROUP BY:([stc].[store_id])	9.102805	False
1576902	--Clustered Index Scan(OBJECT:([store].[production].[stocks].[PK_stoc	10	9	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[production].[stocks].[PK_stocks_E68284D39	7.197171	False
1750000	--Clustered Index Scan(OBJECT:([store].[sales].[orders].[PK_orders_465962	11	6	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[orders].[PK_orders_46596229601F0	7.539393	False
2000000	--Clustered Index Scan(OBJECT:([store].[sales].[customers].[PK_customer CD	12	4	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[customers].[PK_customer CD65C8B	44.59291	False

SQL Query #1 Optimization

```
SELECT
    DISTINCT c.first_name,c.last_name
FROM
    sales.orders AS o INNER JOIN sales.customers AS c on o.customer_id =
c.customer_id
WHERE
o.store_id IN (
    SELECT
        sto.store_id
    FROM
        sales.stores AS sto
WHERE
    sto.store_id IN ( SELECT stc.store_id FROM production.stocks AS stc
                     WHERE stc.quantity > 1000 ))
```

Exection Plan Diagram



Some Analysis

Rows	StmtText	NodeId	Parent	PhysicalOp	LogicalOp	Argument	TotalSubtreeCost	Parallel
798190	SELECT DISTINCT c.first_name,c.last_name FROM Sales.orders AS o INNER JOIN sales.c	1	0				164.1221	False
798190	--Parallelism(Gather Streams)	2	1	Parallelism	Gather Streams		164.1221	True
798190	--Sort(DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC))	3	2	Sort	Distinct Sort	DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC)	162.7942	True
1697910	--Parallelism(Repartition Streams, Hash Partitioning, PARTITION COLUMNS:([c	4	3	Parallelism	Repartition Streams	PARTITION COLUMNS:([c].[first_name], [c].[last_name])	82.89161	True
1697910	--Hash Match(Inner Join, HASH:([o].[customer_id])=[c].[customer_id]), RES	5	4	Hash Match	Inner Join	HASH:([o].[customer_id])=[c].[customer_id]), RESIDUAL:([new_store].[sales].[custom	69.05932	True
1714899	--Parallelism(Repartition Streams, Hash Partitioning, PARTITION COLUMN	6	5	Parallelism	Repartition Streams	PARTITION COLUMNS:([o].[customer_id])	38.04844	True
1714899	--Hash Match(Inner Join, HASH:([sto].[store_id])=[(o].[store_id])	7	6	Hash Match	Inner Join	HASH:([sto].[store_id])=[(o].[store_id])	35.98306	True
950542	--Hash Match(Left Semi Join, HASH:([sto].[store_id])=[(stc].[store_id]	8	7	Hash Match	Left Semi Join	HASH:([sto].[store_id])=[(stc].[store_id])	17.1488	True
978151	--Bitmap(HASH:([sto].[store_id]), DEFINE:([Bitmap1004])	9	8	Bitmap	Bitmap Create	HASH:([sto].[store_id])	2.63793	True
978151	--Parallelism(Repartition Streams, Hash Partitioning, PARTITI	10	9	Parallelism	Repartition Streams	PARTITION COLUMNS:([sto].[store_id])	2.63793	True
978151	--Index Scan(OBJECT:([new_store].[sales].[stores].[_dta_in	11	10	Index Scan	Index Scan	OBJECT:([new_store].[sales].[stores].[_dta_index_stores_9_1061578820_K1] AS [sto])	1.437483	True
1576902	--Parallelism(Repartition Streams, Hash Partitioning, PARTITION	12	8	Parallelism	Repartition Streams	PARTITION COLUMNS:([stc].[store_id])	5.390551	True
1576902	--Index Seek(OBJECT:([new_store].[production].[stocks].[qua	13	12	Index Seek	Index Seek	OBJECT:([new_store].[production].[stocks].[quantity] AS [stc]), SEEK:([stc].[quantity] >	3.47323	True
1750000	--Parallelism(Repartition Streams, Hash Partitioning, PARTITION CO	14	7	Parallelism	Repartition Streams	PARTITION COLUMNS:([o].[store_id])	8.878751	True
1750000	--Index Scan(OBJECT:([new_store].[sales].[orders].[_dta_index_o	15	14	Index Scan	Index Scan	OBJECT:([new_store].[sales].[orders].[_dta_index_orders_9_1173579219_K2_K7_1_3	0.089407	True
2000000	--Parallelism(Repartition Streams, Hash Partitioning, PARTITION COLUMN	16	5	Parallelism	Repartition Streams	PARTITION COLUMNS:([c].[customer_id])	10.77584	True
2000000	--Index Scan(OBJECT:([new_store].[sales].[customers].[_dta_index cu	17	16	Index Scan	Index Scan	OBJECT:([new_store].[sales].[customers].[_dta_index customers_9_1029578706_K1	6.212092	True

What did we make to optimize?

[1] Query Optimization

- Make inner join on customer_id

[2] Adding Indexes

Index In	Columns	Sort Order
orders	Customer_id	ASC
	Store_id	ASC
stocks	quantity	ASC

[3] Parallel Execution

- The system has decided that it could benefit from performing some of the processing in parallel.
- It shouldn't be a cause for concern - generally - and the "duplicate" rows will be eliminated by the later (Gather Streams) operation.

NoSql Query #1

```
db.getCollection("dbo.ordcuststorstoc").aggregate([{$match: {  quantity: {    $gt: 1000  },  first_name: {    $exists: true  },  last_name: {    $exists: true  }  }}, {$group: {  _id: {    first_name: "$first_name",    last_name: "$last_name"  }  }}, {$project: {  first_name: "$_id.first_name",  last_name: "$_id.last_name"  }}, {$sort: {  _id: - 1  }}, {"allowDiskUse": true}])
```

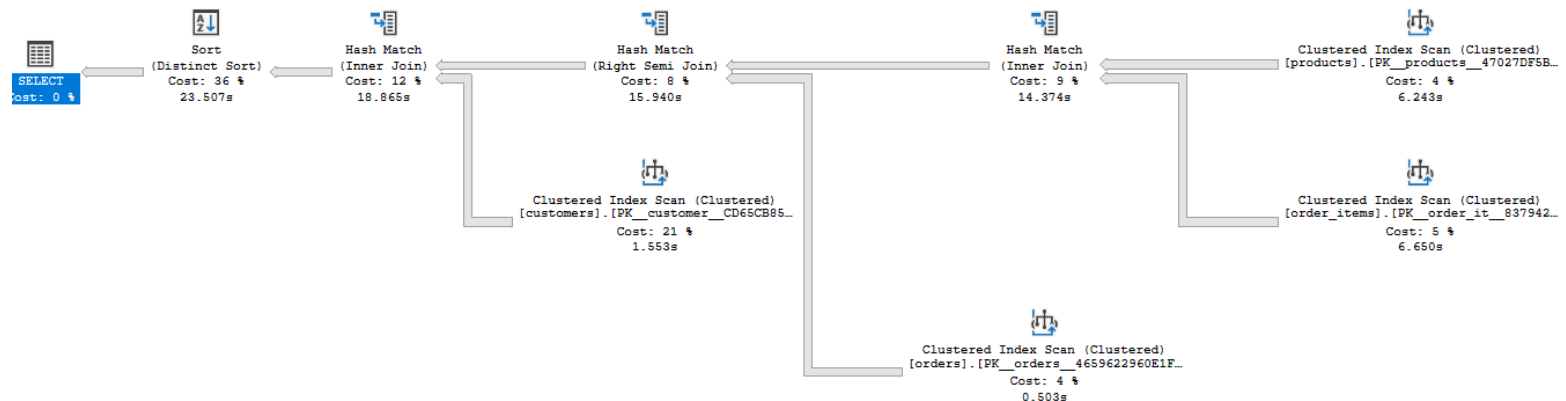
Time Analysis Query #1

Type	Time
SQL (Not Optimized)	00:48.472
Sql (Optimized)	00:22.594
NoSQL	00:67.676

SQL Query #2

```
SELECT
    DISTINCT c.first_name,
    c.last_name
FROM
    sales.orders AS o,
    sales.customers AS c
WHERE
    c.customer_id = o.customer_id
    AND o.shipped_date > '1970-08-11'
    AND o.order_id IN (
        SELECT
            oi.order_id
        FROM
            sales.order_items AS oi
        WHERE
            oi.discount > 10
            AND oi.product_id IN ( SELECT p.product_id FROM production.products AS p WHERE
                p.model_year > 1950 ));
```

Execution Plan Diagram



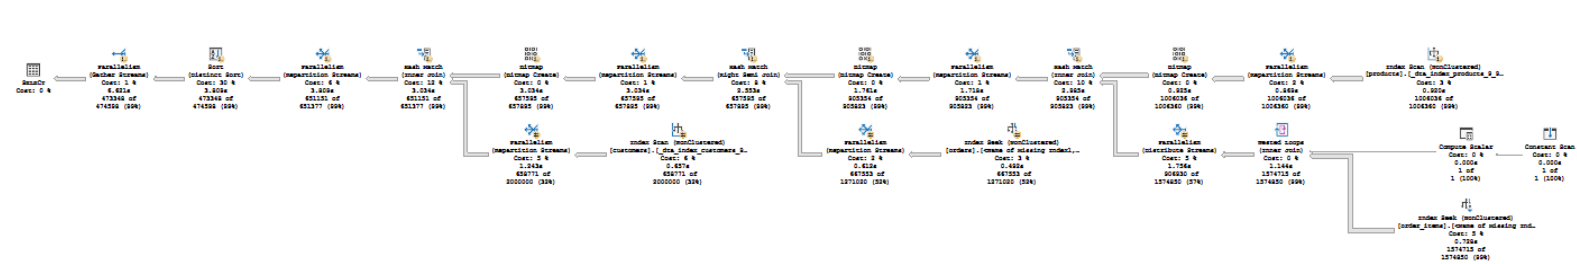
Some Analysis

Rows	StmtText	NodeId	Parent	PhysicalOp	LogicalOp	Argument	TotalSubtreeCost	Parallel
473348	SELECT DISTINCT c.first_name, c.last_name FROM sales.orders AS o, sales.customers AS c	1	0				211.2786	False
473348	--Sort(DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC))	2	1	Sort	Distinct Sort	DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC)	211.2786	False
651151	--Hash Match (Inner Join, HASH:([o].[customer_id])=([c].[customer_id]), RESIDUAL:([o].[order_id])=([c].[order_id]))	3	2	Hash Match	Inner Join	HASH:([o].[customer_id])=([c].[customer_id]), RESIDUAL:([o].[order_id])=([c].[order_id])	135.9496	False
657595	--Hash Match (Right Semi Join, HASH:([o].[order_id])=([o].[order_id]))	4	3	Hash Match	Right Semi Join	HASH:([o].[order_id])=([o].[order_id])	65.59115	False
905354	--Hash Match (Inner Join, HASH:([p].[product_id])=([oi].[product_id]))	6	4	Hash Match	Inner Join	HASH:([p].[product_id])=([oi].[product_id])	40.18327	False
1006036	--Clustered Index Scan (OBJECT:([store].[production].[products].[PK_products__47027DF5B1341017] AS [p]))	7	6	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[production].[products].[PK_products__47027DF5B1341017] AS [p])	9.39643	False
1574715	--Clustered Index Scan (OBJECT:([store].[sales].[order_items].[PK_order_it__837942D42157409F] AS [oi]), WHERE:([oi].[discount] > 10))	8	6	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[order_items].[PK_order_it__837942D42157409F] AS [oi]), WHERE:([oi].[discount] > 10)	11.24458	False
1270917	--Clustered Index Scan (OBJECT:([store].[sales].[orders].[PK_orders__4659622960E1F0CB] AS [o]), WHERE:([o].[shipped_date] > '1970-08-11'))	9	4	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[orders].[PK_orders__4659622960E1F0CB] AS [o]), WHERE:([o].[shipped_date] > '1970-08-11')	7.539393	False
2000000	--Clustered Index Scan (OBJECT:([store].[sales].[customers].[PK_customer_CD65CB8567E6237E] AS [c]))	10	3	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[customers].[PK_customer_CD65CB8567E6237E] AS [c])	44.59291	False

SQL Query #2 Optimization

```
SELECT
    DISTINCT c.first_name,
    c.last_name
FROM
    ( SELECT so.order_id, so.customer_id FROM sales.orders AS so WHERE
so.shipped_date > '1970-08-11' ) AS o
    INNER JOIN sales.customers AS c ON c.customer_id = o.customer_id
WHERE
    o.order_id IN (
        SELECT
            oi.order_id
        FROM
            sales.order_items AS oi
        WHERE
            oi.discount > 10
        AND oi.product_id IN ( SELECT p.product_id FROM production.products AS p WHERE
p.model_year > 1950 ))
```

Exection Plan Diagram



Some Analysis

Rows	StmtText	NodeId	Parent	PhysicalOp	LogicalOp	Argument	TotalSubtreeCost	Parallel
473348	SELECT DISTINCT c.first_name, c.last_name FROM	1	0				95.97579	False
473348	--Parallelism(Gather Streams)	2	1	Parallelism	Gather Streams		95.97579	True
473348	--Sort(DISTINCT ORDER BY:([c].[first_name] ASC, [c]	3	2	Sort	Distinct Sort	DISTINCT ORDER BY:([c].[first_name] ASC, [c].[last_name] ASC)	95.17524	True
651151	--Parallelism(Repartition Streams, Hash Partition	4	3	Parallelism	Repartition Streams	PARTITION COLUMNS:([c].[first_name], [c].[last_name])	66.29818	True
651151	--Hash Match(Inner Join, HASH:([so].[customer	5	4	Hash Match	Inner Join	HASH:([so].[customer_id])=([c].[customer_id]), RESIDUAL:([new	60.92782	True
657595	--Bitmap(HASH:([so].[customer_id]), DEFINE	6	5	Bitmap	Bitmap Create	HASH:([so].[customer_id])	38.30396	True
657595	--Parallelism(Repartition Streams, Hash P	7	6	Parallelism	Repartition Streams	PARTITION COLUMNS:([so].[customer_id])	38.30396	True
657595	--Hash Match(Right Semi Join, HASH:([c	8	7	Hash Match	Right Semi Join	HASH:([oi].[order_id])=([so].[order_id])	37.48721	True
905354	--Bitmap(HASH:([oi].[order_id]), DEF	9	8	Bitmap	Bitmap Create	HASH:([oi].[order_id])	24.92697	True
905354	--Parallelism(Repartition Streams	11	9	Parallelism	Repartition Streams	PARTITION COLUMNS:([oi].[order_id])	24.92697	True
905354	--Hash Match(Inner Join, HASH	12	11	Hash Match	Inner Join	HASH:([p].[product_id])=([oi].[product_id])	23.81318	True
1006036	--Bitmap(HASH:([p].[product	13	12	Bitmap	Bitmap Create	HASH:([p].[product_id])	4.54513	True
1006036	--Parallelism(Repartition	14	13	Parallelism	Repartition Streams	PARTITION COLUMNS:([p].[product_id])	4.54513	True
1006036	--Index Scan(OBJECT:([15	14	Index Scan	Index Scan	OBJECT:([new_store].[production].[products].[_dta_index_pro	2.890889	True
906930	--Parallelism(Distribute Stre	16	12	Parallelism	Distribute Streams	PARTITION COLUMNS:([oi].[product_id]), WHERE:(PROBE([Bitma	10.04797	True
1574715	--Nested Loops(Inner Join	17	16	Nested Loops	Inner Join	OUTER REFERENCES:([Expr1005], [Expr1006], [Expr1004])	5.050431	False
1	--Compute Scalar(DEFIN	18	17	Compute Scalar	Compute Scalar	DEFINE:([([Expr1005],[Expr1006],[Expr1004])=GetRangeWithMism	0	False
1	--Constant Scan	19	18	Constant Scan	Constant Scan		0	False
1574715	--Index Seek(OBJECT:([20	17	Index Seek	Index Seek	OBJECT:([new_store].[sales].[order_items].[<Name of Missing In	5.050431	False
667553	--Parallelism(Repartition Streams, H	21	8	Parallelism	Repartition Streams	PARTITION COLUMNS:([so].[order_id])	4.715209	True
667553	--Index Seek(OBJECT:([new_store	22	21	Index Seek	Index Seek	OBJECT:([new_store].[sales].[orders].[<Name of Missing Index1,	2.681522	True
658771	--Parallelism(Repartition Streams, Hash Part	23	5	Parallelism	Repartition Streams	PARTITION COLUMNS:([c].[customer_id])	10.77584	True
658771	--Index Scan(OBJECT:([new_store].[sales]	24	23	Index Scan	Index Scan	OBJECT:([new_store].[sales].[customers].[_dta_index_customer	6.212092	True

What did we make to optimize?

[1] Query Optimization

- Select first on the order.shipped_date and return only the columns needed later
- Make inner join on customer_id of order and customer table

[2] Adding Indexes

Index In	Columns	Sort Order
products	product_id	ASC
	model_year	ASC
order_items	discount	ASC
orders	shipped_date	ASC

[3] Parallel Execution

- The system has decided that it could benefit from performing some of the processing in parallel.
- It shouldn't be a cause for concern - generally - and the "duplicate" rows will be eliminated by the later (Gather Streams) operation.

NoSql Query #2

```
db.getCollection("dbo.ordcusitempro").aggregate([{$match: {  model_year: {    $gt: 1950  },    shipped_date: {    $gt: new ISODate('1970-08-11')  },    discount: {    $gt: 10  },    first_name: {    $exists: true  },    last_name: {    $exists: true  }  }}, {$group: {  _id: {    first_name: "$first_name",    last_name: "$last_name"  }  }}, {
```



```
$project: {
  first_name: "$_id.first_name",
  last_name: "$id.last_name"
}, {
  $sort: {
    _id: - 1
  }
}], {
  "allowDiskUse": true
})
```

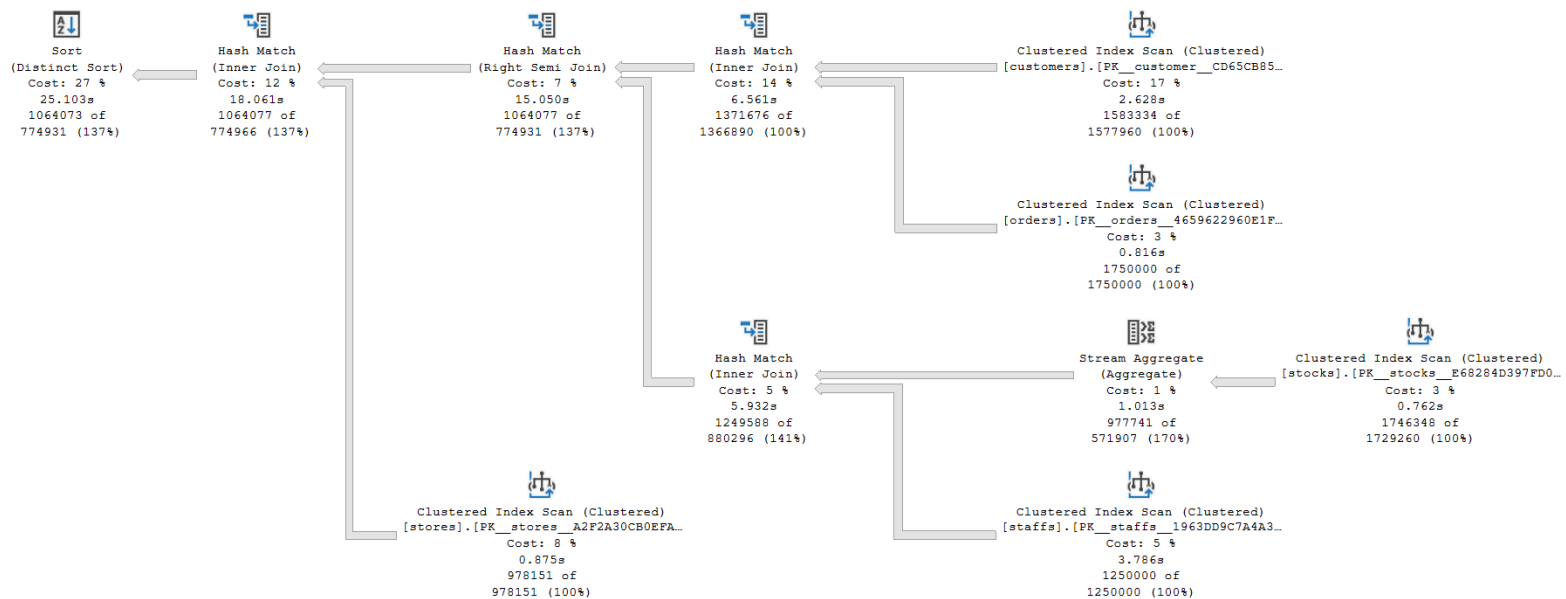
Time Analysis Query #2

Type	Time
SQL (Not Optimized)	00:32.914
Sql (Optimized)	00:14.589
NoSQL	00:58.487

SQL Query #3

```
SELECT
    DISTINCT s.first_name,
    s.last_name,
    sto.store_name
FROM
    sales.staffs AS s,
    sales.stores AS sto
WHERE
    s.store_id = sto.store_id
    AND sto.store_id IN ( SELECT st.store_id FROM production.stocks AS st
WHERE st.quantity > 20 )
    AND s.staff_id IN (
        SELECT
            o.staff_id
        FROM
            sales.orders AS o
        WHERE
            o.staff_id = s.staff_id
            AND o.customer_id IN ( SELECT c.customer_id FROM sales.customers AS c WHERE
                c.zip_code > '2000' ))
```

Execution Plan Diagram



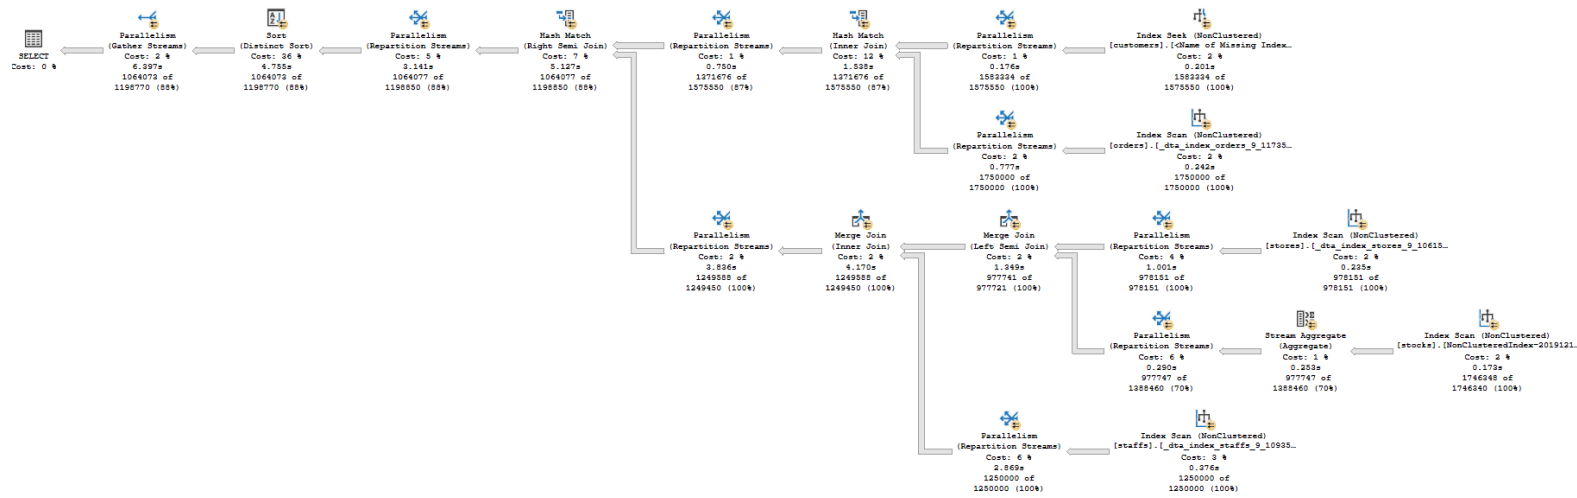
Some Analysis

Row	StmtText	Nodeld	Parent	PhysicalOp	LogicalOp	Argument	TotalSubtreeCost	Parallel
1064073	SELECT DISTINCT s.first_name, s.last_name, s.store_name FROM	1	0				261.471	False
1064073	--Sort(DISTINCT ORDER BY:([s].[first_name] ASC, [s].[last_name] A	2	1	Sort	Distinct Sort	DISTINCT ORDER BY:([s].[first_name] ASC, [s].[last_name] A	261.471	False
1064077	--Hash Match(Inner Join, HASH:([st].[store_id])=([sto].[stor	4	2	Hash Match	Inner Join	HASH:([st].[store_id])=([sto].[store_id])	191.8728	False
1064077	--Hash Match(Right Semi Join, HASH:([o].[staff_id])=([s].[5	4	Hash Match	Right Semi Join	HASH:([o].[staff_id])=([s].[staff_id]), RESIDUAL:([store].[sal	139.2443	False
1371676	--Hash Match(Inner Join, HASH:([c].[customer_id])=([c	7	5	Hash Match	Inner Join	HASH:([c].[customer_id])=([o].[customer_id]), RESIDUAL:([s	88.02207	False
1583334	--Clustered Index Scan(OBJECT:([store].[sales].[cus	8	7	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[customers].[PK_customer_CD65C	44.59291	False
1750000	--Clustered Index Scan(OBJECT:([store].[sales].[ord	9	7	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[orders].[PK_orders_4659622960E	7.539393	False
1249588	--Hash Match(Inner Join, HASH:([st].[store_id])=([s].[s	10	5	Hash Match	Inner Join	HASH:([st].[store_id])=([s].[store_id])	33.21813	False
977741	--Stream Aggregate(GROUP BY:([st].[store_id]))	11	10	Stream Aggregate	Aggregate	GROUP BY:([st].[store_id])	9.187757	False
1746348	--Clustered Index Scan(OBJECT:([store].[product	12	11	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[production].[stocks].[PK_stocks_E68284C	7.197171	False
1250000	--Clustered Index Scan(OBJECT:([store].[sales].[staf	13	10	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[staffs].[PK_staffs_1963D9C7A4A	12.04347	False
978151	--Clustered Index Scan(OBJECT:([store].[sales].[stores].f	14	4	Clustered Index Scan	Clustered Index Scan	OBJECT:([store].[sales].[stores].[PK_stores_A2F2A30CB0E	20.97406	False

SQL Query #3 Optimization

```
SELECT DISTINCT
    s.first_name,
    s.last_name,
    sto.store_name
FROM
    sales.staffs AS s
    INNER JOIN sales.stores AS sto ON sto.store_id = s.store_id
WHERE
    sto.store_id IN ( SELECT st.store_id FROM production.stocks AS st WHERE
st.quantity > 20 )
    AND s.staff_id IN (
SELECT
    o.staff_id
FROM
    sales.orders AS o
WHERE o.customer_id IN ( SELECT c.customer_id FROM sales.customers AS c
WHERE c.zip_code > '2000' ))
```

Exection Plan Diagram



Some Analysis

Rows	StmtText	NodeId	Parent	PhysicalOp	LogicalOp	Argument	TotalSubtreeCost	Parallel
1064073	SELECT DISTINCT [first_name], [last_name], [store_name] FROM [sales]	1	0				156.2407	False
1064073	--Parallelism (Gather Streams)	2	1	Parallelism	Gather Streams		156.2407	True
1064073	--Sort (DISTINCT ORDER BY: ([s].[first_name] ASC, [s].[last_name]	3	2	Sort	Distinct Sort	DISTINCT ORDER BY: ([s].[first_name] ASC, [s].[last_name] ASC, [sto].	153.5176	True
1064077	--Parallelism (Repartition Streams, Hash Partitioning, PARTITIO	4	3	Parallelism	Repartition Streams	PARTITION COLUMNS: ([s].[first_name], [s].[last_name], [sto].	97.95304	True
1064077	--Hash Match (Right Semi Join, HASH: ([o].[staff_id])=([s].[st	5	4	Hash Match	Right Semi Join	HASH: ([o].[staff_id])=([s].[staff_id])	89.58683	True
1371676	--Parallelism (Repartition Streams, Hash Partitioning, PAF	7	5	Parallelism	Repartition Streams	PARTITION COLUMNS: ([o].[staff_id])	32.50592	True
1371676	--Hash Match (Inner Join, HASH: ([c].[customer_id])=([c	8	7	Hash Match	Inner Join	HASH: ([c].[customer_id])=([o].[customer_id]), RESIDUAL: ([new_store	30.58972	True
1583334	--Parallelism (Repartition Streams, Hash Partitioning, PAF	9	8	Parallelism	Repartition Streams	PARTITION COLUMNS: ([c].[customer_id])	5.513364	True
1583334	--Index Seek (OBJECT: ([new_store].[sales].[custo	10	9	Index Seek	Index Seek	OBJECT: ([new_store].[sales].[customers].<Name of Missing Index, sy	3.597162	True
1750000	--Parallelism (Repartition Streams, Hash Partitioning, PAF	11	8	Parallelism	Repartition Streams	PARTITION COLUMNS: ([o].[customer_id])	6.643196	True
1750000	--Index Scan (OBJECT: ([new_store].[sales].[order	12	11	Index Scan	Index Scan	OBJECT: ([new_store].[sales].[orders].[_dta_index_orders_9_1173579	3.853852	True
1249588	--Parallelism (Repartition Streams, Hash Partitioning, PAF	13	5	Parallelism	Repartition Streams	PARTITION COLUMNS: ([s].[staff_id])	45.60593	True
1249588	--Merge Join (Inner Join, MERGE: ([sto].[store_id])=([s]	14	13	Merge Join	Inner Join	MERGE: ([sto].[store_id])=([s].[store_id]), RESIDUAL: ([new_store].	41.96824	True
977741	--Merge Join (Left Semi Join, MERGE: ([sto].[store_id]	15	14	Merge Join	Left Semi Join	MERGE: ([sto].[store_id])=([st].[store_id]), RESIDUAL: ([new_store].	25.37505	True
978151	--Parallelism (Repartition Streams, Hash Partitioning, PAF	16	15	Parallelism	Repartition Streams	PARTITION COLUMNS: ([sto].[store_id]), ORDER BY: ([sto].[store_id] A	9.147228	True
978151	--Index Scan (OBJECT: ([new_store].[sales].[sto	17	16	Index Scan	Index Scan	OBJECT: ([new_store].[sales].[stores].[_dta_index_stores_9_10615788	2.448594	True
977747	--Parallelism (Repartition Streams, Hash Partitioning, PAF	18	15	Parallelism	Repartition Streams	PARTITION COLUMNS: ([st].[store_id]), ORDER BY: ([st].[store_id] ASC	13.69118	True
977747	--Stream Aggregate (GROUP BY: ([st].[store_id])	19	18	Stream Aggregate	Aggregate	GROUP BY: ([st].[store_id])	5.056811	True
1746348	--Index Scan (OBJECT: ([new_store].[product	20	19	Index Scan	Index Scan	OBJECT: ([new_store].[production].[stocks].[NonClusteredIndex-2019	3.853111	True
1250000	--Parallelism (Repartition Streams, Hash Partitioning, PAF	21	14	Parallelism	Repartition Streams	PARTITION COLUMNS: ([s].[store_id]), ORDER BY: ([s].[store_id] ASC)	14.1888	True
1250000	--Index Scan (OBJECT: ([new_store].[sales].[staffs]	22	21	Index Scan	Index Scan	OBJECT: ([new_store].[sales].[staffs].[_dta_index_staffs_9_109357893	4.515148	True

What did we make to optimize?

[1] Query Optimization

- Make inner join on store_id of staff and store table

[2] Adding Indexes

Index In	Columns	Sort Order
staff	staff_id	ASC
	store_id	ASC
stocks	quantity	ASC
	store_id	ASC
orders	customer_id	ASC
	statff_id	ASC
customers	zip_code	ASC

[3] Parallel Execution

- The system has decided that it could benefit from performing some of the processing in parallel.
- It shouldn't be a cause for concern - generally - and the "duplicate" rows will be eliminated by the later (Gather Streams) operation.

NoSql Query #3

```
db.getCollection("dbo.stafstostocordcus").aggregate([{
  $match: {
    quantity: {
      $gt: 20
    },
    zip_code: {
      $gt: "2000"
    },
    store_name: {
      $exists: true
    },
    s_first_name: {
      $exists: true
    },
    s_last_name: {
      $exists: true
    }
  }
}, {
  $group: {
    _id: {
      store_name: "$store_name",
      s_first_name: "$s_first_name",
      s_last_name: "$s_last_name"
    }
  }
}, {
  $project: {
    store_name: "$_id.store_name",
    s_first_name: "$_id.s_first_name",
    s_last_name: "$_id.s_last_name"
  }
}, {
  $sort: {
    _id: - 1
  }
}], {
  "allowDiskUse": true
})
```

Time Analysis Query #3

Type	Time
SQL (Not Optimized)	00:32.914
Sql (Optimized)	00:18.985
NoSQL	00:89.435